


Behavioral Graph Backend Architecture



1. High-level data flow

Browser → Login (session, users.json, credits) → Flask app on Render.

Profile IQ: Submit job → app runs bg.py → Snowflake BEHAVIORALGRAPH → result CSV → S3 dashboard-inputs.

Hedge Fund IQ: Ticker CSVs in S3 aggregated-tickers; metadata in S3 dashboard-inputs/metadata/.

Subscriber IQ: CSVs in S3 svod-acquisition; attribution → SVOD_Churn_Attribution.py → Snowflake PROCESSEDCLICKSTREAM.

Ticket Sales IQ / Tracker: S3 ticket-sales-iq, ticket-sales-tracker; attribution → Snowflake.

Netflix Ranker: app queries Snowflake BEHAVIORALGRAPH.PUBLIC.NETFLIX; cache in S3.

2. Components

Render (Flask / Gunicorn)

app.py: routes, /api/*, auth, job queue, S3/Snowflake clients.

bg.py: behavioral graph pipeline (universe scan, boosting, ESPN consistency); runs per Profile IQ job.

Attribution scripts: Talent_Theater_Attribution, Ticket_Sales_Attribution, SVOD_Churn_Attribution, etc.

AWS S3 (us-east-2)

dashboard-inputs: profile CSVs, metadata/, purgatory/, system/jobs_status.json, Netflix cache.

svod-acquisition, ticket-sales-iq, ticket-sales-tracker, aggregated-tickers.

Snowflake

BEHAVIORALGRAPH (BEHAVIORGRAPH6X): clickstream, Profile IQ, NETFLIX table.

PROCESSEDCLICKSTREAM: attribution (Ticket, SVOD, Talent).

Data ingestion: lordata_part1/part2 → CSV; send_to_snow.py → Snowflake PUT + COPY.

External: OpenAI (insights, persona, deck); Gmail API (emails).

3. Profile IQ job flow

1. User submits via /api/submit. 2. app creates job_id, status in S3 jobs_status.json.

3. run_analysis(): bg.connect_snowflake(), perform_full_universe_scan(), run_full_pipeline().

4. bg.py reads clickstream, applies pipeline, writes CSV. 5. app uploads CSV to S3, updates cache.

6. User opens profile (data from S3 via presigned URL or API).

4. IQ products

Profile IQ: Run analysis or open existing profile; demographics, behavioral, AI insights.

Hedge Fund IQ: Ticker → linked profiles; SEC actuals, accuracy.

Subscriber IQ / Ticket Sales IQ / Tracker: list and data from S3; run attribution jobs.

Netflix Ranker: by date, show, season, episode from Snowflake or S3 cache.

5. Auth & admin

Session auth; users.json; credits per run. Admin: user CRUD, Gmail, purgatory, cache refresh, ticker metadata.